

# lab7\_pin\_page 实验说明

本实验是为了验证和学习如何去 pin 一个 page，也就是熟悉 `get_user_pages()` 和 `get_user_pages_fast()` 的使用。这两个接口函数在内核代码或者驱动代码中很常见。

## 基本实验步骤

1. 进入 `runninglinuxkernel_4.0/rlk_lab/rlk_basic/chapter_7_mm/lab7_pin_page` 目录。

```
# export ARCH=arm
# export CROSS_COMPILE=arm-linux-gnueabi-
# make BASEINCLUDE=/home/figo/work/runninglinuxkernel/runninglinuxkernel_4.0
```

这里 `BASEINCLUDE` 指定你当前 `runninglinuxkernel_4.0` 的目录路径。

编译 test 测试 app。

```
# arm-linux-gnueabi-gcc test_ok.c -o test_ok
```

然后把 `mydevdemo-pin-page.ko` 和 `test_ok` 拷贝到 `runninglinuxkernel_4.0/kmodules` 目录下。

运行如下脚本启动 Qemu。

```
#cd runninglinuxkernel_4.0
# sh run.sh arm32 #启动虚拟机
```

在 Qemu 虚拟机:

```
#cd /mnt
# insmod mydevdemo-pin-page.ko
#./test_ok
```

```

/mnt # insmod mydevdemo-pin-page.ko
succeeded register char device: my_demo_dev
/mnt # ./test_ok
demodrv_open: major=10, minor=58
driver max buffer size=4096
demodrv_read_write: len=4096, npage=1
pin 1 pages from user done
demodrv_read_write: write user buffer 4096 bytes done
demodrv_write: write nbytes=4096 done at pos=0
demodrv_read_write: len=4096, npage=1
pin 1 pages from user done
demodrv_read_write: read user buffer 4096 bytes done
demodrv_read: read nbytes=4096 done at pos=0
data modify and compare succussful
/mnt # █

```

可以看到测试成功了！

## 进阶思考

笨叔在实验里，设置一个进阶思考的问题。

代码里有一个 test\_issue.c 文件，最大的不同就是使用 malloc 来分配 user buffer，而不是通过 mmap 来分配的匿名页面。

```

34
35     read_buffer = malloc(len);
36     if (!read_buffer)
37         goto open_fail;
38
39     write_buffer = malloc(len);
40     if (!write_buffer)
41         goto buffer_fail;
42

```

运行 test\_issue 程序，得到如下结果：

```

/mnt # ./test_issue
demodrv_open: major=10, minor=58
driver max buffer size=4096
demodrv_read_write: len=4096, npage=1
pin 1 pages from user done
demodrv_read_write: write user buffer 4096 bytes done
demodrv_write: write nbytes=4096 done at pos=0
demodrv_read_write: len=4096, npage=1
pin 1 pages from user done
demodrv_read_write: read user buffer 4096 bytes done
demodrv_read: read nbytes=4096 done at pos=0
buffer compare fail

```

程序的最后结果是“buffer compare fail”，说明程序运行有问题，也就是 read buffer 和 write buffer 数据不对，究竟怎么回事呢？

遇到这种问题，我们最简单最粗暴也是最有效的调试办法就是把 buffer 打印出来看看。在内核里，可以使用 `print_hex_dump_bytes()` 函数。比如在 `demodrv_read_write()` 函数的第 61,67,72 行添加打印语句。

```
59     for (i = 0; i < npages; i++) {
60         kmap_addr = kmap(pages[i]);
61         //print_hex_dump_bytes("kmap:", DUMP_PREFIX_OFFSET, kmap_addr, PAGE_SIZE);
62         size = min_t(size_t, PAGE_SIZE, len);
63         switch(rw) {
64             case MYDEMO_READ:
65                 memcpy(kmap_addr, dev_buf + PAGE_SIZE * i,
66                        size);
67                 //print_hex_dump_bytes("read:", DUMP_PREFIX_OFFSET, kmap_addr, size);
68                 break;
69             case MYDEMO_WRITE:
70                 memcpy(dev_buf + PAGE_SIZE * i, kmap_addr,
71                        size);
72                 //print_hex_dump_bytes("write:", DUMP_PREFIX_OFFSET, dev_buf + PAGE_SIZE * i, size);
73                 break;
74             default:
75                 break;
76         }
```

打印结果如下：

```
pin 1 pages from user done
kmap:00000000: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000010: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000020: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000030: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000040: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000050: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000060: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000070: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000080: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000090: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000a0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000b0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000c0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000d0: 00 00 00 00 09 10 00 00 55 55 55 55 55 55 55 55 .....UUUUUUUU
kmap:000000e0: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUUUU
kmap:000000f0: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUUUU
kmap:00000100: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUUUU
kmap:00000110: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUUUU
kmap:00000120: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUUUU
kmap:00000130: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUUUU
kmap:00000140: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUUUU
```

发现 pin 的 page，开头的数字都是 0，而不是 0x55，这是为什么呢？理论上这个 page 应该全部都是 0x55 才对。

这是为什么呢？

如果大家对这个问题感兴趣，可以关注笨叔的第一季旗舰篇视频，笨叔会在视频中和大家详细解答。

shop115683645.taobao.com

配套视频

旗舰篇

第1季

内存管理

奔跑吧Linux内核

基于Linux 4.x 内核源代码问题分解2

奔跑吧Linux社区

旗舰篇一次订阅，持续更新

规划中

第二季

进程管理和调度 / 中断 / 锁（已出）

第三季

虚拟化

第四季

Linux 内核和应用开发调试必杀技

第五季

红帽系列

shop115683645.taobao.com

奔跑吧Linux社区

Linux视频课程

微信公众号：奔跑吧 linux 社区

1. >一键订阅，持续更新

2. >最有深度和广度的 Linux 视频

3. >手把手解读 Linux 内核代码

4. >紧跟 Linux 开源社区技术热点

5. >笨叔叔的 VIP 私密群答疑

6. >图书 + 视频，全新学习模式

微店：



扫码识别

淘宝店: <https://shop115683645.taobao.com/>

微信公众号:

