

lab2_mutex 实验说明

在第 5 章的虚拟 FIFO 设备中，我们并没有考虑多个进程同时访问设备驱动的情况，请使用互斥锁对虚拟 FIFO 设备驱动程序进行并发保护

基本实验步骤

1. 进入 `rlk_lab/rlk_basic/chapter_9/lab2` 目录。

```
# export ARCH=arm
# export CROSS_COMPILE=arm-linux-gnueabi-
# make BASEINCLUDE=/home/figo/work/runninglinuxkernel/runninglinuxkernel_4.0
```

这里 `BASEINCLUDE` 指定你当前 `runninglinuxkernel_4.0` 的目录路径。

编译 `test` 测试 app。

```
# arm-linux-gnueabi-gcc test.c -o test
```

然后把 `mydemo_fasync.ko` 和 `test` 拷贝到 `runninglinuxkernel_4.0/kmodules` 目录下面。

运行如下脚本启动 Qemu。

```
#cd runninglinuxkernel_4.0
# sh run.sh arm32 #启动虚拟机
```

在 Qemu 虚拟机:

```
#cd /mnt
# insmod mydemo_fasync.ko
```

```

/mnt # insmod mydemo_fasync.ko
my_class mydemo:252:0: create device: 252:0
mydemo_fifo=ee098f5c
my_class mydemo:252:1: create device: 252:1
mydemo_fifo=ee0981dc
my_class mydemo:252:2: create device: 252:2
mydemo_fifo=ee09829c
my_class mydemo:252:3: create device: 252:3
mydemo_fifo=ee09835c
my_class mydemo:252:4: create device: 252:4
mydemo_fifo=ee09841c
my_class mydemo:252:5: create device: 252:5
mydemo_fifo=ee09889c
my_class mydemo:252:6: create device: 252:6
mydemo_fifo=ee098e9c
my_class mydemo:252:7: create device: 252:7
mydemo_fifo=ee09859c
succeeded register char device: mydemo_dev

```

你会看到创建了 8 个设备。你可以到 `/sys/class/my_class/` 目录下面看到这些设备。

```

/mnt # cd /sys/class/my_class/
/sys/class/my_class # ls
mydemo:252:0  mydemo:252:2  mydemo:252:4  mydemo:252:6
mydemo:252:1  mydemo:252:3  mydemo:252:5  mydemo:252:7
/sys/class/my_class #

```

我们可以看到创建了主设备号为 252 的设备。我们再来看一下 `/dev/` 目录。

```

/sys/class/my_class # ls -l /dev
total 0
crw-rw----  1 0      0          14,   4 Feb  1 09:46 audio
crw-rw----  1 0      0           5,   1 Feb  1 09:46 console
crw-rw----  1 0      0        10,  63 Feb  1 09:46 cpu_dma_latency
crw-rw----  1 0      0        14,   3 Feb  1 09:46 dsp
crw-rw----  1 0      0        29,   0 Feb  1 09:46 fb0
crw-rw----  1 0      0        29,   1 Feb  1 09:46 fb1
crw-rw----  1 0      0         1,   7 Feb  1 09:46 full
crw-rw----  1 0      0       10, 183 Feb  1 09:46 hwrng
drwxr-xr-x  2 0      0       120 Feb  1 09:46 input
crw-rw----  1 0      0         1,   2 Feb  1 09:46 kmem
crw-rw----  1 0      0         1,  11 Feb  1 09:46 kmsg
crw-rw----  1 0      0         1,   1 Feb  1 09:46 mem
crw-rw----  1 0      0       10,  60 Feb  1 09:46 memory_bandwidth
crw-rw----  1 0      0        14,   0 Feb  1 09:46 mixer
crw-rw----  1 0      0        90,   0 Feb  1 09:46 mtd0

```

发现并没有主设备为 252 的设备。

所以我们需要手工创建一个设备用来 test app。

```
#mknod /dev/mydemo0 c 252 1
```

接下来跑我们的 test 程序：

```
# ./test &    #这里让 test 程序在后台跑
```

```
/mnt # ./test &
/mnt # my_class mydemo:252:1: demodrv_open: major=252, minor=1, device=mydemo_dev1
my_class mydemo:252:1: demodrv_fasync send SIGIO
```

然后使用 echo 命令来往/dev/mydemo0 这个设备写入字符串。

```
/mnt # echo "i am study linux now" > /dev/mydemo0
my_class mydemo:252:1: demodrv_open: major=252, minor=1, device=mydemo_dev1
demodrv_write kill fasync
my_class mydemo:252:1: demodrv_write:mydemo_dev1 pid=700, actual_write =21, ppos=0, ret=0
FIFO is not empty
my_class mydemo:252:1: demodrv_read:mydemo_dev1, pid=772, actual_readed=21, pos=0
i am study linux now
```

可以看到从 demodrv_read()函数把刚才写入的字符串已经读到用户空间了。

进阶思考

笨叔在实验里，设置一个进阶思考的问题。

代码里有一个 test_issue.c 文件，最大的不同就是使用 malloc 来分配 user buffer，而不是通过 mmap 来分配的匿名页面。

```
34
35     read_buffer = malloc(len);
36     if (!read_buffer)
37         goto open_fail;
38
39     write_buffer = malloc(len);
40     if (!write_buffer)
41         goto buffer_fail;
42
```

运行 test_issue 程序，得到如下结果：

```
/mnt # ./test_issue
demodrv_open: major=10, minor=58
driver max buffer size=4096
demodrv_read_write: len=4096, npage=1
pin 1 pages from user done
demodrv_read_write: write user buffer 4096 bytes done
demodrv_write: write nbytes=4096 done at pos=0
demodrv_read_write: len=4096, npage=1
pin 1 pages from user done
demodrv_read_write: read user buffer 4096 bytes done
demodrv_read: read nbytes=4096 done at pos=0
buffer compare fail
```

程序的最后结果是“buffer compare fail”，说明程序运行有问题，也就是 read buffer 和 write buffer 数据不对，究竟怎么回事呢？

遇到这种问题，我们最简单最粗暴也是最有效的调试办法就是把 buffer 打印出来看看。在内

核里,可以使用 `print_hex_dump_bytes()`函数。比如在 `demodrv_read_write()`函数的第 61,67,72 行添加打印语句。

```
59     for (i = 0; i < npages; i++) {
60         kmap_addr = kmap(pages[i]);
61         //print_hex_dump_bytes("kmap:", DUMP_PREFIX_OFFSET, kmap_addr, PAGE_SIZE);
62         size = min_t(size_t, PAGE_SIZE, len);
63         switch(rw) {
64             case MYDEMO_READ:
65                 memcpy(kmap_addr, dev_buf + PAGE_SIZE * i,
66                        size);
67                 //print_hex_dump_bytes("read:", DUMP_PREFIX_OFFSET, kmap_addr, size);
68                 break;
69             case MYDEMO_WRITE:
70                 memcpy(dev_buf + PAGE_SIZE*i, kmap_addr,
71                        size);
72                 //print_hex_dump_bytes("write:", DUMP_PREFIX_OFFSET, dev_buf + PAGE_SIZE*i, size);
73                 break;
74             default:
75                 break;
76         }
```

打印结果如下:

```
pin 1 pages from user done
kmap:00000000: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000010: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000020: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000030: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000040: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000050: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000060: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000070: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000080: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:00000090: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000a0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000b0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000c0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 .....
kmap:000000d0: 00 00 00 00 09 10 00 00 55 55 55 55 55 55 55 55 .....UUUUUUUU
kmap:000000e0: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUU
kmap:000000f0: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUU
kmap:00000100: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUU
kmap:00000110: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUU
kmap:00000120: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUU
kmap:00000130: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUU
kmap:00000140: 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 55 UUUUUUUUUUUUUUUUU
```

发现 `pin` 的 `page`, 开头的数据都是 0, 而不是 0x55, 这是为什么呢? 理论上这个 `page` 应该全部都是 0x55 才对。

这是为什么呢?

如果大家对这个问题感兴趣, 可以关注笨叔的第一季旗舰篇视频, 笨叔会在视频中和大家详细解答。

shop115683645.taobao.com

配套视频 **旗舰篇**

第**1**季
内存管理



**奔跑吧
Linux 社区**

旗舰篇一次订阅，持续更新

规划中

第二季	进程管理和调度 / 中断 / 锁（已出）
第三季	虚拟化
第四季	Linux 内核和应用开发调试必杀技
第五季	红帽系列

shop115683645.taobao.com

Linux 视频课程



微信公众号：奔跑吧 linux 社区

1. > 一键订阅，持续更新
2. > 最有深度和广度的 Linux 视频
3. > 手把手解读 Linux 内核代码
4. > 紧跟 Linux 开源社区技术热点
5. > 笨叔叔的 VIP 私密群答疑
6. > 图书 + 视频，全新学习模式

微店：



扫码识别

淘宝店: <https://shop115683645.taobao.com/>

微信公众号:

